

Import

In [17]:

```
import pandas as pd
import numpy as np
from keras.models import Sequential
from keras.layers import Dense
from keras.wrappers.scikit_learn import KerasClassifier
from keras.utils import np_utils
from sklearn import linear_model
from sklearn.model_selection import cross_val_score, KFold, GridSearchCV, RandomizedSearchCV
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score, mean_squared_error, r2_score
from sklearn.pipeline import Pipeline
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
from keras.optimizers import Adam
import warnings
warnings.filterwarnings('ignore')
```

Fuel Consumption Dataset - Regression

Data load and clean

In [18]:

```
fuel = pd.read_csv('https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/ML0101ENv3/labs/FuelConsumptionCo2.csv')
```

In [19]:

```
fuel.head()
```

Out[19]:

	MODELYEAR	MAKE	MODEL	VEHICLECLASS	ENGINE SIZE	CYLINDERS	TRANSMISSION
0	2014	ACURA	ILX	COMPACT	2.0	4	
1	2014	ACURA	ILX	COMPACT	2.4	4	
2	2014	ACURA	ILX HYBRID	COMPACT	1.5	4	
3	2014	ACURA	MDX 4WD	SUV - SMALL	3.5	6	
4	2014	ACURA	RDX AWD	SUV - SMALL	3.5	6	

In [20]:

```
print(fuel.info())
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1067 entries, 0 to 1066
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	MODELYEAR	1067 non-null	int64
1	MAKE	1067 non-null	object
2	MODEL	1067 non-null	object
3	VEHICLECLASS	1067 non-null	object
4	ENGINE SIZE	1067 non-null	float64
5	CYLINDERS	1067 non-null	int64
6	TRANSMISSION	1067 non-null	object
7	FUELTYPE	1067 non-null	object
8	FUELCONSUMPTION_CITY	1067 non-null	float64
9	FUELCONSUMPTION_HWY	1067 non-null	float64
10	FUELCONSUMPTION_COMB	1067 non-null	float64
11	FUELCONSUMPTION_COMB_MPG	1067 non-null	int64
12	CO2EMISSIONS	1067 non-null	int64

```
dtypes: float64(4), int64(4), object(5)
```

```
memory usage: 108.5+ KB
```

```
None
```

In [21]:

```
#fuel['CYLINDERS'] = fuel['CYLINDERS'].astype('object')
print(fuel.MODELYEAR.unique())
fuel.drop('MODELYEAR', axis = 1, inplace = True)
fuel.drop('MODEL', axis = 1, inplace = True)
fuel.drop('MAKE', axis = 1, inplace = True)
fuel.head()
```

[2014]

Out[21]:

	VEHICLECLASS	ENGINE SIZE	CYLINDERS	TRANSMISSION	FUELTYPE	FUELCONSUMPTION
0	COMPACT	2.0	4	AS5	Z	
1	COMPACT	2.4	4	M6	Z	
2	COMPACT	1.5	4	AV7	Z	
3	SUV - SMALL	3.5	6	AS6	Z	
4	SUV - SMALL	3.5	6	AS6	Z	

In [22]:

```
fuel = pd.get_dummies(fuel)
fuel.head()

# OR

# for i in fuel.describe(include = np.object).columns:
#     fuel[i] = fuel[i].astype('category')
#     fuel[i] = fuel[i].cat.codes
#     fuel[i] = fuel[i].astype('category')
```

Out[22]:

	ENGINE SIZE	CYLINDERS	FUELCONSUMPTION_CITY	FUELCONSUMPTION_Hwy	FUELCONSUMPTION_COMB
0	2.0	4	9.9	6.7	8.1
1	2.4	4	11.2	7.7	9.4
2	1.5	4	6.0	5.8	5.9
3	3.5	6	12.7	9.1	10.9
4	3.5	6	12.1	8.7	10.4

5 rows × 49 columns

In [23]:

fuel.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1067 entries, 0 to 1066

Data columns (total 49 columns):

#	Column	Non-Null Count	Dtype
0	ENGINE SIZE	1067 non-null	float64
1	CYLINDERS	1067 non-null	int64
2	FUEL CONSUMPTION_CITY	1067 non-null	float64
3	FUEL CONSUMPTION_HWY	1067 non-null	float64
4	FUEL CONSUMPTION_COMB	1067 non-null	float64
5	FUEL CONSUMPTION_COMB_MPG	1067 non-null	int64
6	CO2 EMISSIONS	1067 non-null	int64
7	VEHICLE CLASS_COMPACT	1067 non-null	uint8
8	VEHICLE CLASS_FULL-SIZE	1067 non-null	uint8
9	VEHICLE CLASS_MID-SIZE	1067 non-null	uint8
10	VEHICLE CLASS_MINICOMPACT	1067 non-null	uint8
11	VEHICLE CLASS_MINIVAN	1067 non-null	uint8
12	VEHICLE CLASS_PICKUP TRUCK - SMALL	1067 non-null	uint8
13	VEHICLE CLASS_PICKUP TRUCK - STANDARD	1067 non-null	uint8
14	VEHICLE CLASS_SPECIAL PURPOSE VEHICLE	1067 non-null	uint8
15	VEHICLE CLASS_STATION WAGON - MID-SIZE	1067 non-null	uint8
16	VEHICLE CLASS_STATION WAGON - SMALL	1067 non-null	uint8
17	VEHICLE CLASS_SUBCOMPACT	1067 non-null	uint8
18	VEHICLE CLASS_SUV - SMALL	1067 non-null	uint8
19	VEHICLE CLASS_SUV - STANDARD	1067 non-null	uint8
20	VEHICLE CLASS_TWO-SEATER	1067 non-null	uint8
21	VEHICLE CLASS_VAN - CARGO	1067 non-null	uint8
22	VEHICLE CLASS_VAN - PASSENGER	1067 non-null	uint8
23	TRANSMISSION_A4	1067 non-null	uint8
24	TRANSMISSION_A5	1067 non-null	uint8
25	TRANSMISSION_A6	1067 non-null	uint8
26	TRANSMISSION_A7	1067 non-null	uint8
27	TRANSMISSION_A8	1067 non-null	uint8
28	TRANSMISSION_A9	1067 non-null	uint8
29	TRANSMISSION_AM5	1067 non-null	uint8
30	TRANSMISSION_AM6	1067 non-null	uint8
31	TRANSMISSION_AM7	1067 non-null	uint8
32	TRANSMISSION_AS4	1067 non-null	uint8
33	TRANSMISSION_AS5	1067 non-null	uint8
34	TRANSMISSION_AS6	1067 non-null	uint8
35	TRANSMISSION_AS7	1067 non-null	uint8
36	TRANSMISSION_AS8	1067 non-null	uint8
37	TRANSMISSION_AS9	1067 non-null	uint8
38	TRANSMISSION_AV	1067 non-null	uint8
39	TRANSMISSION_AV6	1067 non-null	uint8

```
40 TRANSMISSION_AV7          1067 non-null  uint8
41 TRANSMISSION_AV8          1067 non-null  uint8
42 TRANSMISSION_M5           1067 non-null  uint8
43 TRANSMISSION_M6           1067 non-null  uint8
44 TRANSMISSION_M7           1067 non-null  uint8
45 FUELTYPE_D                 1067 non-null  uint8
46 FUELTYPE_E                 1067 non-null  uint8
47 FUELTYPE_X                 1067 non-null  uint8
48 FUELTYPE_Z                 1067 non-null  uint8
dtypes: float64(4), int64(3), uint8(42)
memory usage: 102.2 KB
```

In [24]:

```
fuel.head()
```

Out[24]:

	ENGINE_SIZE	CYLINDERS	FUEL_CONSUMPTION_CITY	FUEL_CONSUMPTION_HWY	FUEL_COST_PER_GAL
0	2.0	4	9.9	6.7	12.0
1	2.4	4	11.2	7.7	10.0
2	1.5	4	6.0	5.8	15.0
3	3.5	6	12.7	9.1	7.0
4	3.5	6	12.1	8.7	7.4

5 rows × 49 columns

Regression

In [25]:

```
x = fuel.drop('CO2EMISSIONS', axis = 1)
y = fuel['CO2EMISSIONS']

x_train, x_test, y_train, y_test = train_test_split(x, y, train_size = 0.6)
```

In [26]:

```
linear = linear_model.LinearRegression()
ridge = linear_model.Ridge()
lasso= linear_model.Lasso()
elastic = linear_model.ElasticNet()
lasso_lars = linear_model.LassoLars()
bayes_ridge = linear_model.BayesianRidge()
logistics = linear_model.LogisticRegression()
sgd = linear_model.SGDClassifier()
passagg = linear_model.PassiveAggressiveClassifier()
ridgecv = linear_model.RidgeClassifierCV()
ridgeclass = linear_model.RidgeClassifier()

models_churn = [linear, ridge, lasso, elastic, lasso_lars, bayes_ridge, logist
ics, sgd, passagg, ridgecv, ridgeclass]
```

In [27]:

```
def get_cv_scores(model):
    scores = cross_val_score(model, x_train, y_train, cv=5, scoring='neg_root_
mean_squared_error')
    print('CV Mean: ', np.mean(scores))
    print('STD: ', np.std(scores))
    print('\n')
```

In [28]:

```
for i in models_churn:
    print(i)
    get_cv_scores(i)
```

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, nor
malize=False)
CV Mean:  -5.568129523825883
STD:  0.736552577166622
```

```
Ridge(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=None,
      normalize=False, random_state=None, solver='auto', tol=0.001
)
CV Mean:  -5.609066774025289
STD:  0.8041810228548972
```

```
Lasso(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=1000,
      normalize=False, positive=False, precompute=False, random_st
ate=None,
      selection='cyclic', tol=0.0001, warm_start=False)
CV Mean:  -8.91487029464611
STD:  1.5701638550563086
```

```
ElasticNet(alpha=1.0, copy_X=True, fit_intercept=True, l1_ratio=0.5,
           max_iter=1000, normalize=False, positive=False, precompute=False,
           random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
CV Mean: -20.372459824583984
STD: 2.761703817318036
```

```
LassoLars(alpha=1.0, copy_X=True, eps=2.220446049250313e-16, fit_intercept=True,
           fit_path=True, max_iter=500, normalize=True, positive=False,
           precompute='auto', verbose=False)
CV Mean: -32.217937653789605
STD: 3.7430626640724554
```

```
BayesianRidge(alpha_1=1e-06, alpha_2=1e-06, alpha_init=None,
               compute_score=False, copy_X=True, fit_intercept=True,
               lambda_1=1e-06, lambda_2=1e-06, lambda_init=None, n_iter=300,
               normalize=False, tol=0.001, verbose=False)
CV Mean: -5.540264290478805
STD: 0.7369043882596081
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=1000,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001,
                    verbose=0, warm_start=False)
CV Mean: -20.31535769981978
STD: 0.8534805377822611
```

```
SGDClassifier(alpha=0.0001, average=False, class_weight=None,
               early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
               l1_ratio=0.15, learning_rate='optimal', loss='hinge',
               max_iter=1000, n_iter_no_change=5, n_jobs=None, penalty='l2',
               power_t=0.5, random_state=None, shuffle=True, tol=0.001,
               validation_fraction=0.1, verbose=0, warm_start=False)
CV Mean: -56.025583409900534
```

STD: 3.6584154466486605

```
PassiveAggressiveClassifier(C=1.0, average=False, class_weight=None,
                             early_stopping=False, fit_intercept=True,
                             loss='hinge', max_iter=1000, n_iter_no
                             _change=5,
                             n_jobs=None, random_state=None, shuffle=True,
                             tol=0.001, validation_fraction=0.1, verbose=0,
                             warm_start=False)
```

CV Mean: -45.81636097300676

STD: 6.679853455813455

```
RidgeClassifierCV(alphas=array([ 0.1, 1. , 10. ]), class_weight=None, cv=None,
                   fit_intercept=True, normalize=False, scoring=None,
                   store_cv_values=False)
```

CV Mean: -40.06360226773207

STD: 2.5586785325580603

```
RidgeClassifier(alpha=1.0, class_weight=None, copy_X=True, fit_intercept=True,
                 max_iter=None, normalize=False, random_state=None,
                 solver='auto', tol=0.001)
```

CV Mean: -39.33511788091995

STD: 4.115527158716331

In [29]:

```
alpha = [0.5, 1, 1.5, 0.01, 2.5, 0.0001, 10, 100, 0.35]
#solver = ['auto', 'svd']

param_grid = dict(alpha = alpha)#, solver = solver)
```

In [30]:

```
grids = GridSearchCV(estimator = lasso_lars, param_grid = param_grid, scoring = 'r2', cv = 10)
grid_result = grids.fit(x_train, y_train)
```


In [31]:

```
rnds = RandomizedSearchCV(estimator = ridge, param_distributions = param_grid,
scoring='r2', cv = 10)
rnds_result = rnds.fit(x_train, y_train)
```

In [32]:

```
print(grid_result.best_params_)
print(rnds_result.best_params_)
```

```
{'alpha': 0.0001}
{'alpha': 0.35}
```

Best Model

In [34]:

```
print(grid_result.best_score_)
print(rnds_result.best_score_)
```

```
0.9916810501375295
0.9917607214483969
```

Iris Dataset - Classification

In [35]:

```

from sklearn.datasets import load_iris
iris = load_iris()
data1 = pd.DataFrame(data= np.c_[iris['data'], iris['target']],
                      columns= iris['feature_names'] + ['target'])
data1

```

Out[35]:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
0	5.1	3.5	1.4	0.2	0.0
1	4.9	3.0	1.4	0.2	0.0
2	4.7	3.2	1.3	0.2	0.0
3	4.6	3.1	1.5	0.2	0.0
4	5.0	3.6	1.4	0.2	0.0
...
145	6.7	3.0	5.2	2.3	2.0
146	6.3	2.5	5.0	1.9	2.0
147	6.5	3.0	5.2	2.0	2.0
148	6.2	3.4	5.4	2.3	2.0
149	5.9	3.0	5.1	1.8	2.0

150 rows × 5 columns

In [36]:

```

x_iris = data1.drop('target', axis = 1)
y_iris = data1['target']
x1_train, x1_test, y1_train, y1_test = train_test_split(x_iris, y_iris, train_
size = 0.7)
print(x1_train.shape)
print(x1_test.shape)
print(y1_train.shape)
print(y1_test.shape)

```

(105, 4)

(45, 4)

(105,)

(45,)

Model Building

In [37]:

```
def get_cv_scores1(model):
    scores = cross_val_score(model, x1_train, y1_train, cv=5, scoring='accuracy')
    print('CV Mean Score: ', np.mean(scores))
    print('STD: ', np.std(scores))
    print('\n')

models_log = [logistics, sgd, passagg, ridgecv, ridgeclass]
```

In [38]:

```
for i in models_log:
    print(i)
    get_cv_scores1(i)
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001,
                    verbose=0,
                    warm_start=False)
CV Mean Score: 0.9428571428571428
STD: 0.05553287518900288
```

```
SGDClassifier(alpha=0.0001, average=False, class_weight=None,
              early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
              l1_ratio=0.15, learning_rate='optimal', loss='hinge',
              max_iter=1000, n_iter_no_change=5, n_jobs=None, penalty='l2',
              power_t=0.5, random_state=None, shuffle=True, tol=0.001,
              validation_fraction=0.1, verbose=0, warm_start=False)
CV Mean Score: 0.780952380952381
STD: 0.06459361888690732
```

```
PassiveAggressiveClassifier(C=1.0, average=False, class_weight=None,
                            early_stopping=False, fit_intercept=True,
                            loss='hinge', max_iter=1000, n_iter_no_change=5,
                            n_jobs=None, random_state=None, shuffle=True,
                            tol=0.001, validation_fraction=0.1, ve
```

```

rbose=0,
                                warm_start=False)
CV Mean Score:  0.9142857142857144
STD:  0.06317380553057904

```

```

RidgeClassifierCV(alphas=array([ 0.1,  1. , 10. ]), class_weight=None,
                  cv=None,
                  fit_intercept=True, normalize=False, scoring=None,
                  store_cv_values=False)
CV Mean Score:  0.838095238095238
STD:  0.07126966450997983

```

```

RidgeClassifier(alpha=1.0, class_weight=None, copy_X=True, fit_intercept=True,
               max_iter=None, normalize=False, random_state=None,
               solver='auto', tol=0.001)
CV Mean Score:  0.838095238095238
STD:  0.07126966450997983

```

In [39]:

```

grid_list = []
random_list = []
for i in models_log:
    grids = GridSearchCV(estimator = lasso_lars, param_grid = param_grid, scoring = 'r2', cv = 10)
    grid_result = grids.fit(x_train, y_train)
    grid_list.append(grid_result.best_score_)
    rnds = RandomizedSearchCV(estimator = ridge, param_distributions = param_grid, scoring='r2', cv = 10)
    rnds_result = rnds.fit(x_train, y_train)
    random_list.append(rnds_result.best_score_)

```

Best model

In [40]:

```

print(max(grid_list))
print(max(random_list))

```

```

0.9916810501375295
0.9917607214483969

```

In []: